

Q-Framework Development & Validation

Whitepaper v1.0 RC1

Compute-State Virtualization and Adaptive Flow Control for Heterogeneous AI Infrastructure

October 2026

| This paper documents a running system.

Q-Framework is an experimental heterogeneous AI compute architecture developed to investigate whether long-context generative workloads can remain recoverable and controllable when their practical working set approaches or exceeds the comfortable physical-memory envelope of a consumer GPU.

The current implementation includes explicit state lineage, checkpoint/resume behavior, CPU/RAM-backed state handling, runtime control signals, static-state sources, multi-node continuation, and candidate working-set virtualization inside model-forward execution.

The public paper reports observable architecture, development milestones, validation methods and current results. Proprietary scheduling, state-selection, adaptive-control and production data-plane algorithms are intentionally excluded.

Design principles

1. Logical state and physical residency are different concepts.
2. Memory pressure should be observable and actionable.
3. Compute state should survive temporary resource pressure.
4. A GPU is an execution resource, not the entire machine.
5. Claims must be tied to runtime evidence.

Current evidence

Demonstrated evidence includes deterministic state/hash verification across GPU continuation paths, real GPU→CPU/RAM→GPU rescue lineages, candidate QmRNA execution inside H3 model-forward blocks, exact FFN token chunking, and CPU-only global streamed-attention reference tests matching full attention within floating-point tolerance.

The complete target of a final production-quality 15-second full-length H3 task on a single 22GB GPU through the newest full-length working-set virtualization path is still under validation and is not presented as completed fact.

Disclosure boundary

This publication intentionally omits production state-selection policy, checkpoint packing, restore sequencing, pressure scoring, Qsearch ranking, scheduler weights, lease/fencing rules and the production streamed-QKV data-plane.